

## CSC\_4SL05\_TP/SLR206B: project pseudocode

Our multi-writer multi-reader atomic register implementation (a variant of ABD [1]) is described in Algorithm 1.

To write a new value  $v$ , a writer increments its local sequence number  $r$ , sends a *read request*  $[?, r]$  to all the processes, waits for the responses of the type  $[v', t', r]$  from a *majority* ( $> n/2$ ) of the processes (we also say a *quorum*). Then it picks up the highest timestamp  $t_{max}$  in the set of received responses, computes its new timestamp as  $t = t_{max} + 1$  and sends a *write request*  $[v, t]$  to all the processes in the system. The operation completes as soon as responses of the type  $[ack, v, t]$  are received from a quorum.

To read the register value, a reader increments its sequence number  $r$  and sends a *read request*  $[?, r]$  to every all the processes. As soon as responses of the type  $[v', t', r]$  are received from a quorum, the operation selects the value  $v_m$  equipped with the highest timestamp  $t_m$ . If there are multiple different values written with timestamp  $t_m$ , the largest such value is chosen. To make sure that subsequent read operations will not “miss” the returned value, the reader then sends a write request  $[v_m, t_m]$ , waits until a quorum confirms it by sending  $[ack, v_m, t_m]$ , and only then returns  $v$ .

Notice that, as multiple values can be written with the same timestamp, the value is attached to the acknowledgement sent to a write request. This is needed to ensure that the acknowledgement is indeed sent to *this* write request and *not* to an earlier write request with the same timestamp but a different value (this could happen if the operation is concurrent with multiple writes using the same timestamp). An alternative solution is to maintain an additional sequence number attached to every write request a process issues.

## References

- [1] H. Attiya, A. Bar-Noy, and D. Dolev. Sharing memory robustly in message passing systems. *J. ACM*, 42(2):124–142, Jan. 1995.

---

**Algorithm 1** Multi-writer multi-reader atomic register: pseudocode

---

```
1: Local variables:  
2:   localValue = 0 // locally stored value, initially 0  
3:   localTS = 0 // locally stored timestamp, initially 0  
4:   t = 0 // timestamp  
5:   r = 0 // sequence number (the number of issued read requests)  
  
6: Upon write(v):  
7:   r +=  
8:   send [?, r] to all // read request  
9:   wait until received {[v', t', r]} from a majority  
10:  let t' be the highest received timestamp  
11:  t = t' + 1 // new timestamp  
12:  send [v, t] to all // write request  
13:  wait until received [ack, v, t] from a majority  
14:  return ok  
  
15: Upon read(v):  
16:   r +=  
17:   send [?, r] to all // read request  
18:   wait until received {[v', t', r]} from a majority  
19:   let  $v_m$  be the largest value with the highest timestamp  $t_m$   
20:   send [ $v_m, t_m$ ] to all // write request  
21:   wait until received [ack,  $v_m, t_m$ ] from a majority  
22:   return  $v_m$   
  
23: Upon received [ $v', t'$ ] from  $p_j$ :  
24:   if  $t' > localTS$  or ( $t' = localTS$  and  $v' > localValue$ ) then  
25:     localValue :=  $v'$   
26:     localTS :=  $t'$   
27:     send [ack,  $v', t'$ ] to  $p_j$   
  
28: Upon received [?,  $r'$ ] from  $p_j$ :  
29:   send [localValue, localTS,  $r'$ ] to  $p_j$ 
```

---