

CSC_4SL05_TP/SLR206B: Project Robust Key-Value Store

The goal of this project is to get an initial experience in designing a fault-tolerant distributed system. Here we focus on the popular key-value store application.

1 Sequential specification

The state of a key-value store is a set of key-value pairs of the form (k, v) , where k is an integer and v is a value in a given *value set* (assume that values are also integers). The initial state is an empty set. The system exports two operations:

- $\text{put}(k, v)$ sets the value with key k to v (overwriting the old value if it is already in the set).
- $\text{get}(k)$ returns the value of key k (the default value \perp is returned if the key is not at the system).

2 Concurrent environment

The goal of the project is to give a key-value store implementation for the following environment:

- We have N asynchronous processes. Every process has a distinct *identifier*. The identifiers are publicly known.
- Every two processes can communicate via a reliable asynchronous point-to-point channel.
- Up to $f < N/2$ of the processes are subject to crash failures: a faulty process prematurely stops taking steps of its algorithm. A process that never crashes is called correct.

The implementation should ensure that in its every execution, the following conditions are met:

Safety The corresponding history is linearizable with respect to the sequential specification above;

Liveness Every operation invoked by a correct process eventually returns.

Hint: You can treat the system as a set of multi-writer multi-reader atomic registers (indexed by keys). Just as in the ABD algorithm (check the class on distributed storage or the lecture notes), for each key k , every process maintains a local copy of the value, equipped with a *timestamp*. To write a new value v , we need to make sure that at least a majority ($N - f$) of the processes store v or a newer value in their local copies. To read the value, we need to contact a majority of the processes and return the most recent received value.

Of course, the ABD algorithm only implements a *single-writer regular* register (in the multi-reader case), while we want to implement a *multi-writer atomic* one. So the writer (the process executing a *put* operation) needs to be careful to choose a timestamp for the message it writes to be higher than the timestamps of preceding writes. Respectively, to prevent the new-old inversion, the reader (the process executing a *get* operation) needs to ensure that the returned value (or a newer one) is stored at a majority of the processes.

This might require one extra round of message exchange between the process performing an operation and a *quorum* (a majority) of other processes.

3 Prerequisites

The project assumes a basic knowledge of Java. Get familiarized with the Java version of AKKA, an actor-based programming model <https://akka.io/docs/>. Check basic constructions in to see how to create an actor, and make the actors communicate.

Check <https://github.com/remisharrock/SLR210Patterns> for sample AKKA patterns which you might want to use.

4 Formalities

The project is pursued in **teams of three students** each (*exceptionally*, can also be two or four).

The implemented system should be provided with a short report describing how the system operates and containing correctness arguments.

5 Implementation

The implementation should extend the basic construction creating a system of a given size and ensure all-to-all connectivity. More precisely, in the main class, create N actors (processes), and pass references of all N processes to each of them. Use the name **Process** for the process class.

In the **Process** class create methods for executing operations *put* and *get*. **For simplicity, implement just the system for just a single key. Recall that this is equivalent to implementing a single atomic register.**

To test the implementation and measure its performance, use the following procedure.

The **main** method selects f processes at random (e.g., using the **shuffle** method from `java.collections`) and sends each of them a special *crash* message. If a process receives a *crash* message it enters the *silent* mode, not reacting to any future event.

For every remaining process, the **main** method sends a special *launch* message. Once process i receives a *launch* message, it sequentially performs:

- M *put* operations, with parameters $k = 1$ and $v = i, N + i, 2 * N + i, \dots, M * N + i$, and
- M *get* operations, with parameters $k = 1$.

Make sure that every process performs at most one operation at a time (remember that we require every exported history to be *well-formed*).

Use the **LoggingAdapter** class to log both the invocation and the response of each operation a process performs together with its timing.

Perform the experiment for $N = 3, 10, 100$ (with $f = 1, 4$, and 49 , respectively) and $M = 3, 10, 100$ (nine instances). For the instance $N = 3$ and $M = 3$, check that the resulting execution is linearizable. Also, for each instance, measure the *latency*, i.e., the total computation time.

6 Report

Prepare a short report (up to 15 pages), preferably in English (can also be written in French if English does not feel comfortable). The report should contain:

- A high level description of the system;
- A sketch of a proof of correctness (please argue that both safety and liveness hold);
- A report on performance analysis.

The report and the java source of the implementation should be submitted by via moodle (archived in one zip file) by the deadline set on the course page.

7 Presentation

If the time permits, we'll have a short (up to 7 mins) discussion of the reports. Please be prepared to answer the questions on the code, as well as on the outcomes of your experiments.